

Our course agenda

- Introduction and overview
- NumPy: Basic data handling with Numpy arrays
- Pandas
 - Exploratory data analysis
 - Data consolidation
 - Data cleaning
- Data visualization using Matplotlib and Seaborn
- Interacting with APIs
- Interacting with SQL databases
- Version Control with Git and GitHub
- Advanced Python

Python foundations



Operators

Functions

Control flow and iterators

Programming concepts & paradigms

See also Precourse Programming

Tooling

Installation

Visual Studio Code

Jupyter Notebooks


Packages

Virtual Environments

Git and Github



Python










Pandas in the Data Science Process



Pandas Series and DataFrame

Pandas Series

index

| 0 | "apples" |
|---|-----------|
| 1 | "bananas" |
| 2 | "cherries |

Pandas DataFrame

columns

index

| | fruit | amount |
|---|-----------|--------|
| 0 | "apples" | 3 |
| 1 | "bananas" | 2 |
| 2 | "cherries | 5 |

```
pd.Series(["apples", "bananas", "cherries"])
```

Attributes of a DataFrame

| | fruit | amount |
|---|-----------|--------|
| 0 | "apples" | 3 |
| 1 | "bananas" | 2 |
| 2 | "cherries | 5 |

Index

| | _ | | Truit | aate | amount | price | |
|------------------------------------|-----|-----------|------------|------------|--------|-------|--|
| _ | [| 0 | "apples" | 2024-03-01 | 3 | 2.99 | |
| | | 1 | "bananas" | 2024-03-01 | 2 | 1.99 | |
| | | 2 | "cherries | 2024-03-01 | 5 | 3.49 | |
| 16 / [6 17 | | | | | | | |
| <pre>df.set_index(["fruit"])</pre> | | fruit | date | amount | price | | |
| | L.[| "apples" | 2024-03-01 | 3 | 2.99 | | |
| | | "bananas" | 2024-03-01 | 2 | 1.99 | | |
| | | "cherries | 2024-03-01 | 5 | 3.49 | | |

df.reset_index()

- ▶ It can be useful to use an index based on one or multiple columns of the DataFrame:
 - Fast selection of data
 - Intuitive and fast merge of multiple DataFrames
 - Efficient operations on time series data
 - Some visualization packages work naturally with index columns

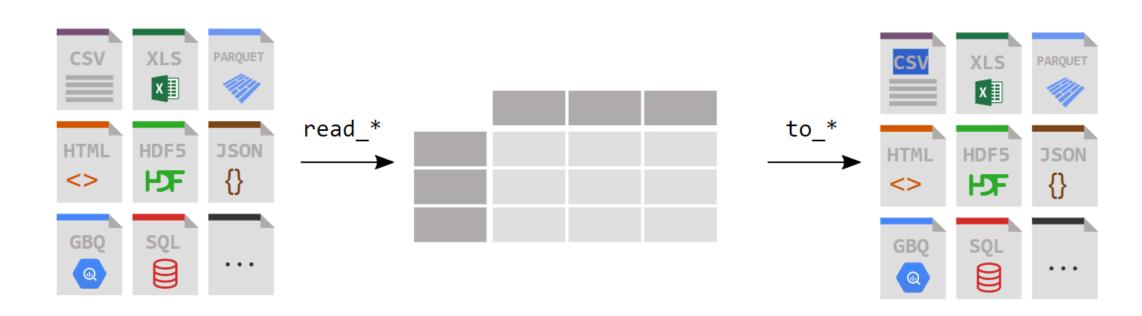
Data Input and Output

Typically, we read in data from files

- pd.read_csv(filename)
- pd.read_excel(filename)
- pd.read_sql()
- ...

Or write data back into some files

- df.to_csv(filename)
- df.to_excel(filename)
- df.to_sql()
- ...







.iloc accessor

Selecting rows and columns based on INTEGER LOCATION

- .iloc[row_indices, col_indices]
- We specifiy row and column indices in same way as for Numpy arrays
 - start:stop:step
 - Zero-based
 - Negative indexing allowed

| fruit | date | amount | price |
|-----------|------------|--------|-------|
| "apples" | 2024-03-01 | 3 | 2.99 |
| "bananas" | 2024-03-01 | 2 | 1.99 |
| "cherries | 2024-03-01 | 5 | 3.49 |

| fruit | amount | price |
|-----------|--------|-------|
| "apples" | 3 | 2.99 |
| "bananas" | 2 | 1.99 |

loc accessors

Selecting rows and columns based on LABELS

| fruit | date | amount | price |
|-----------|------------|--------|-------|
| "apples" | 2024-03-01 | 3 | 2.99 |
| "bananas" | 2024-03-01 | 2 | 1.99 |
| "cherries | 2024-03-01 | 5 | 3.49 |

Index column

df.loc["apples", "price"]

| fruit | amount | price |
|-----------|--------|-------|
| "apples" | 3 | 2.99 |
| "bananas" | 2 | 1.99 |

Selecting only columns

| | fruit | date | amount | price |
|---|-----------|------------|--------|-------|
| 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |

| | date | price |
|---|------------|-------|
| 0 | 2024-03-01 | 2.99 |
| 1 | 2024-03-01 | 1.99 |
| 2 | 2024-03-01 | 3.49 |
| 3 | 2024-03-01 | 2.49 |

- → returns DataFrame with 1 column
- ▶ df["price"] or df.price → returns Series

Selecting rows using boolean mask

Conditional Expression that evaluates to a Pandas Series of length equal to number of rows in DataFrame

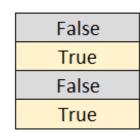
| | | fruit | date | amount | price |
|---|---|-----------|------------|--------|-------|
| | 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| - | 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| | 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| | 3 | "dates" | 2024-03-01 | 8 | 2.49 |

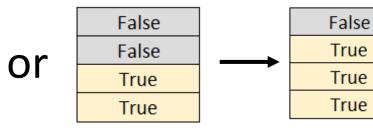
| False | |
|-------|--|
| True | |
| False | |
| True | |
| | |

| | | fruit | date | amount | price |
|---|---|-----------|------------|--------|-------|
| ١ | 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| | 3 | "dates" | 2024-03-01 | 8 | 2.49 |

Multiple conditions: or operator "|"

| | | fruit | date | amount | price |
|---|---|-----------|------------|--------|-------|
| | 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| - | 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| | 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| | 3 | "dates" | 2024-03-01 | 8 | 2.49 |

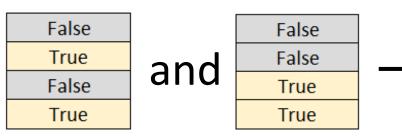




| | fruit | date | amount | price |
|---|-----------|------------|--------|-------|
| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |

Multiple conditions: and operator "&"

| _ | | fruit | date | amount | price |
|----|---|-----------|------------|--------|-------|
| | 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| ·[| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| | 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| | 3 | "dates" | 2024-03-01 | 8 | 2.49 |



df[(df.price < 2.5) & (df.amount > 4)]

| | fruit | date | amount | price |
|---|---------|------------|--------|-------|
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |

False

False

False

True





Derive new columns from existing ones

| | fruit | amount | price |
|---|-----------|--------|-------|
| 1 | "bananas" | 2 | 1.99 |
| 2 | "cherries | 5 | 3.49 |
| 3 | "dates" | 8 | 2.49 |

```
df["sales"] = df["amount"] * df["price"]

df = df.assign(sales = df["amount"] * df["price"])
```

| | | fruit | amount | price | sales |
|------------------|---|-----------|--------|-------|-------|
| $\left[\right]$ | 1 | "bananas" | 2 | 1.99 | 3.98 |
| | 2 | "cherries | 5 | 3.49 | 17.45 |
| | 3 | "dates" | 8 | 2.49 | 19.92 |

- Calculations are performed elementwise
- No need for a loop
- To overwrite an existing column with new values we can use the same syntax
- The assign method creates a new DataFrame, while the first approach modifies the existing DataFrame

Sorting

- > sort_values: sort according to a column / list of multiple columns
- sort_index: sort according to an index / multiple indices
- Argument ascending=False for inverse sorting

| | fruit | date | amount | price |
|---|-----------|------------|--------|-------|
| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |
| 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| 2 | "cherries | 2024-03-01 | 5 | 3.49 |

df.sort_values("price")

| | fruit | date | amount | price |
|---|-----------|------------|--------|-------|
| 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| 2 | "cherries | 2024-03-01 | 5 | 3.49 |
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |

df.sort_index()





Aggregation Functions

- Statistical methods for Series and DataFrames
 - count: number of non-missing observations
 - sum, mean, min, max
 - std: standard deviation
 - var: variance
 - cumsum: cumulative sum
 - •
- ▶ Most operations return a <u>single aggregated value</u> (sum, mean, ...) per column

| | fruit | date | amount | price |
|---|-----------|------------|--------|-------|
| 1 | "bananas" | 2024-03-01 | 2 | 1.99 |
| 3 | "dates" | 2024-03-01 | 8 | 2.49 |
| 0 | "apples" | 2024-03-01 | 3 | 2.99 |
| 2 | "cherries | 2024-03-01 | 5 | 3.49 |

Useful aggregation arguments

- axis: axis of the DataFrame along which the statistic is computed (default: 0)
- numeric_only: calculate statistics only for numeric columns (default: False)
- **skipna**: ignore NaN (missing values) (default: True)

| | amount | price | sales | | | |
|-----------------|--------|-------|-------|-----------------|-----------|-------|
| "bananas" | 2 | 1.99 | 3.98 | df.mean(axis=1) | "bananas" | 2.66 |
| "cherries | 5 | 3.49 | 17.45 | | "cherries | 8.65 |
| "dates" | 8 | 2.49 | 19.92 | | "dates" | 10.14 |
| df.mean(axis=0) | | | | | | |

| amount | price | sales |
|--------|-------|-------|
| 5 | 2.66 | 13.78 |

agg method

The agg method is useful for more complex aggregation scenarios. It accepts:

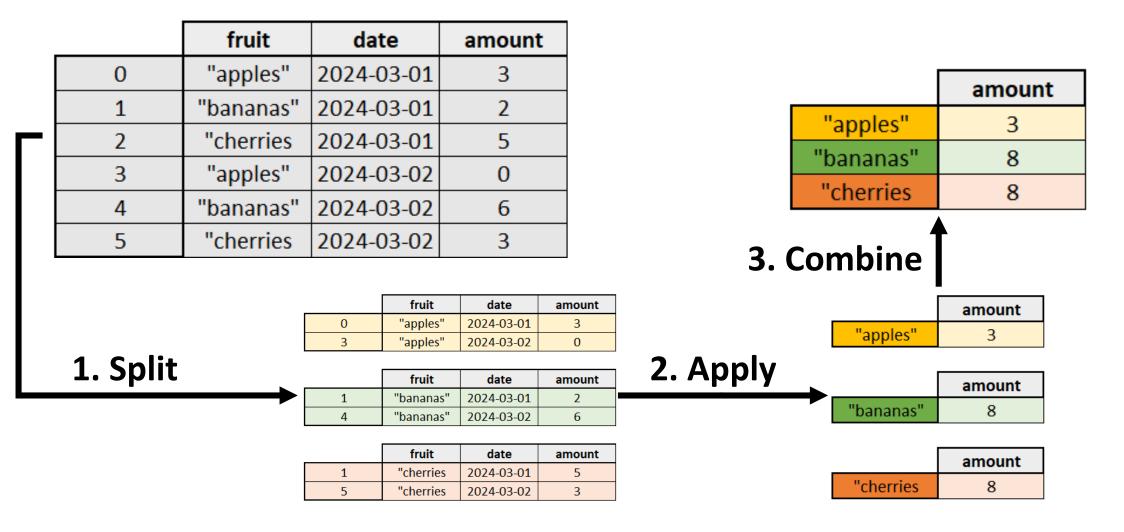
- Custom aggregation function
- String function name: 'mean', 'max'
- <u>List</u> of multiple aggregation functions: [np.mean, 'max']
- Dictionary with column <-> aggregation function mapping

| | amount | price |
|-----------|--------|-------|
| "bananas" | 2 | 1.99 |
| "cherries | 5 | 3.49 |
| "dates" | 8 | 2.49 |
| | | |
| | | |





Grouped Aggregation: split - apply - combine



df.groupby("fruit").amount.sum()

Method chaining

```
df.groupby("fruit").amount.sum()
```

- We can arrange a sequence of Pandas operations in a chain for the purpose of
 - Readability
 - Consisness
 - Maintainability
- Optionally, we can break long lines

```
(df
...groupby('fruit')
...amount
...sum()
)
```